

# Overfitting/Underfitting in Training

1

Dr. Yi Long (Neal)

Most contents (text or images) of course slides are from the following textbook  
Provost, Foster, and Tom Fawcett. Data Science for Business: What you need to  
know about data mining and data-analytic thinking. " O'Reilly Media, Inc.", 2013



# Outline

- Hinge Loss & SVM
- Nonlinear Models & Other Extensions
- Overfitting & Holdout Evaluation
- Quiz

# Margin Defined by Hinge Loss


- Hinge Loss defines two hyperplanes are parallel to the decision boundary, and a margin between the two parallel hyperplanes

$$\text{hinge\_loss}(f, (\mathbf{x}, y)) = \max(0, 1 - yf(\mathbf{x}))$$




# Hinge Loss and Clean Margin

- By minimizing hinge loss is actually trying to find decision boundary with a relative clean margin



# Revisit of Quiz-6 Q2



# Distance between Points/Lines

## Distance from point to straight line

- [https://en.wikipedia.org/wiki/Distance\\_from\\_a\\_point\\_to\\_a\\_line](https://en.wikipedia.org/wiki/Distance_from_a_point_to_a_line)

$$\text{distance}(ax + by + c = 0, (x_0, y_0)) = \frac{|ax_0 + by_0 + c|}{\sqrt{a^2 + b^2}}.$$

## Distance between straight lines

- [https://en.wikipedia.org/wiki/Distance\\_between\\_two\\_straight\\_lines](https://en.wikipedia.org/wiki/Distance_between_two_straight_lines)

When the lines are given by

$$ax + by + c_1 = 0$$

$$ax + by + c_2 = 0,$$

the distance between them can be expressed as

$$d = \frac{|c_2 - c_1|}{\sqrt{a^2 + b^2}}.$$

## Distance from point to hyperplane

$$f(\mathbf{x}) = w_0 + w_1 * x_1 + w_2 * x_2 + \dots = 0$$

$$f(\mathbf{x}) = w_0 + \mathbf{w}^T \cdot \mathbf{x} = 0 = 0$$

$$\text{distance}(f(\mathbf{x}) = 0, \mathbf{x}_0) = \frac{f(\mathbf{x}_0)}{\|\mathbf{w}\|_2}$$

## Distance between hyperplanes

$$f_1(\mathbf{x}) = c_1 + \mathbf{w}^T \cdot \mathbf{x} = 0$$

$$f_2(\mathbf{x}) = c_2 + \mathbf{w}^T \cdot \mathbf{x} = 0$$

$$\text{distance}(f_1(\mathbf{x}) = 0, f_2(\mathbf{x}) = 0) = \frac{|c_2 - c_1|}{\|\mathbf{w}\|_2}$$

# Intuition of Support Vector Machine (SVM)

- SVM tries to find some trade-off between **a fat margin** and **a low total loss**

- ✓ A fat margin is a penalty for complexity (restrict the searching of  $\mathbf{w}$ )

- Fat margin bar = Bar with large margin width between two parallel boundaries

- ✓ Margin width between boundaries is  $\frac{2}{\|\mathbf{w}\|_2}$


- ✓ Can maximize margin by minimizing  $\|\mathbf{w}\|_2$

- ✓ Therefore, SVM tries to optimize

$$\min_{\mathbf{w}} \text{HingeLoss}(\mathbf{w}) + C \cdot \|\mathbf{w}\|_2$$

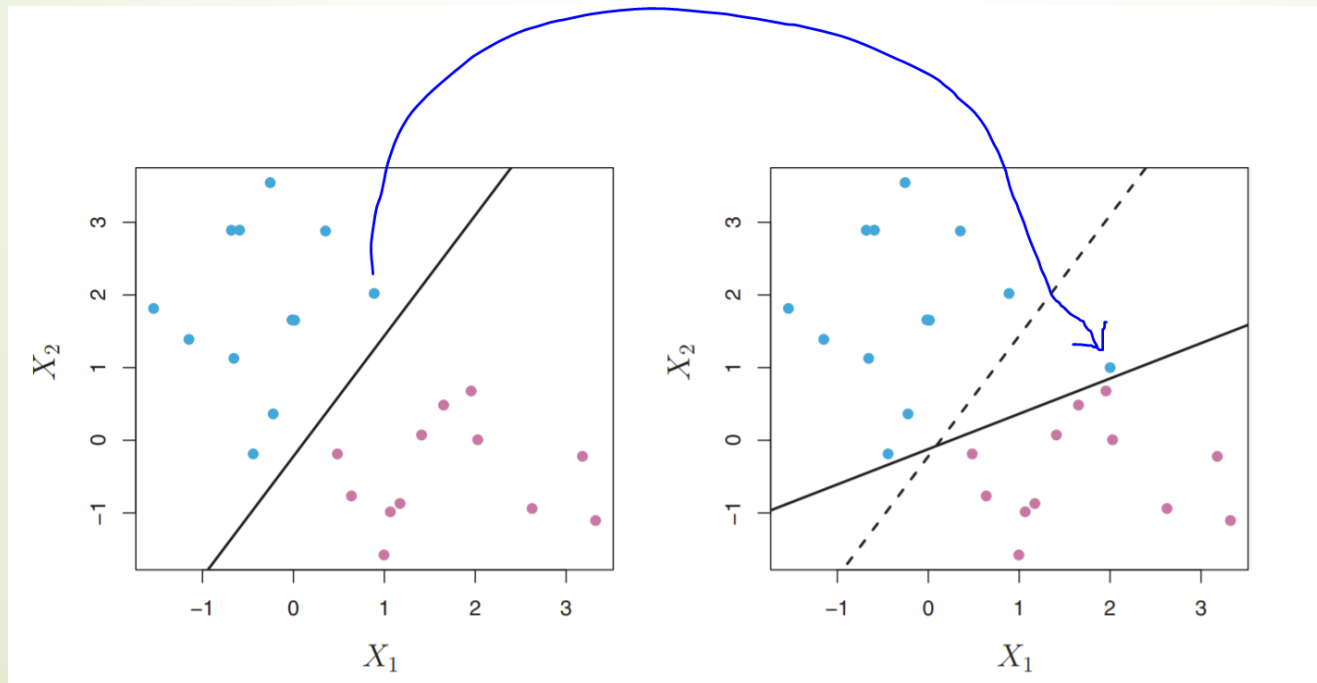
- $C$  is the coefficient for the trade-off

- Trade-off is everywhere





# Intuition of Fat Margin

- ▶ We may want a classifier that may **misclassify a few** observations but with **fat margin** for:
  - ▶ Robustness for individual observation.
  - ▶ Better classification of **most** of the training observation (**rather than all**)





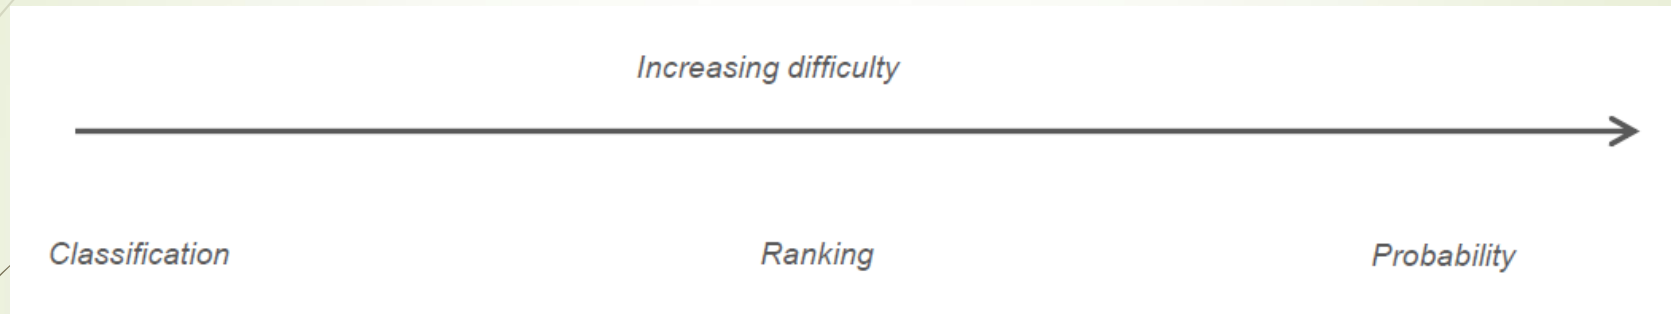
# SVM vs. Logistic Regression



# Outline

- Hinge Loss & SVM
- Nonlinear Models & Other Extensions
- Overfitting & Holdout Evaluation
- Quiz


# Classification/Ranking/Probability Estimation



- With estimated probability, we can compute the expected loss of fraud detection/churn prediction
  - ✓ We can further classify or rank samples easily
- Ranking can make classification results more useful
  - ✓ Locate limited budget to award/retain those users to be churned

# Distance of Linear Discriminant Functions


- ▶ When a sample near the decision boundary, we would be most uncertain about its class
- ▶ When it goes further away from the decision boundary, we would expect the higher likelihood of predicted class based on sign of linear discriminant functions
- ▶ **Distance** provides us free **ranking measure** of likelihood of correct estimation
- ▶ Logistic regression **directly** interpret the distance as class probability



# Nonlinear Decision Boundaries

- Linear functions can actually represent nonlinear decision boundaries in higher dimension
  - if we extend the original feature by more complex features, and then we actually extend linear functions to nonlinear decision boundaries accordingly

Add Sepal width<sup>2</sup> to the original feature



# Nonlinear Models-Example

- ▶ We can extend 2-d feature vector of each sample (both training and new data) by adding quadratic combinations to a 5-d feature vector

$$(x_1, x_2) \rightarrow (x_1, x_2, x_1^2, x_2^2, x_1x_2) \rightarrow (x'_1, x'_2, x'_3, x'_4, x'_5)$$


- ▶ Accordingly, the linear discriminant function in 5-d space changes to

$$f(\mathbf{x}) = w_0 + w_1x'_1 + w_2x'_2 + w_3x'_3 + w_4x'_4 + w_5x'_5$$

- ▶ Then we can create a hyperplane in the 5-d space as follows to separate the data as good as the nonlinear function in 2-d space with


$$w_0 = -0.54, w_1 = -0.8, w_2 = 0$$

$$w_3 = 1, w_4 = 1, w_5 = 0$$




# Other Nonlinear Models

- Nonlinear support vector machine with a “polynomial kernel” consider “higher-order” combinations of the original features
  - ‘linear’, ‘poly’, ‘rbf’, ‘sigmoid’, ‘precomputed’ kernel in SVM
- Neural network as a “stack” of models
  - ✓ On the bottom of the stack are the original features
  - ✓ Each layer in the stack applies a simple model to the outputs of the previous layer
  - ✓ Middle layer tries to learn meaningful representation of original feature
- Might fit data too well (overfitting)





Regression





# Representation Learning


Deep neural networks learn hierarchical feature representations






# Binary to Multi-class classification

- We have studied binary classification extensively, but multi-class classification



# One-versus-rest

- We can utilize binary classification model to do multi-class classification by training multiple binary classifier, and choose the predicted label with highest votes/probability



# Recommended Resources

- Visualization: <https://www.w3resource.com/graphics/matplotlib/>
- Data Collection:
  - Selenium: <https://selenium-python.readthedocs.io/>
  - Chrome: F12 is your friend
  - Fiddler: <https://www.telerik.com/blogs/how-to-capture-android-traffic-with-fiddler>
  - Python Scrapy: <https://docs.scrapy.org/en/latest/intro/tutorial.html>

# Outline

- Hinge Loss & SVM
- Nonlinear Models & Other Extensions
- Overfitting & Holdout Evaluation
- Quiz

# Logistic Regression vs. Decision Tree

- Predict whether is cell (with image) are breast cancer or not

Attribute name	Description
RADIUS	<i>Mean of distances from center to points on the perimeter</i>
TEXTURE	<i>Standard deviation of grayscale values</i>
PERIMETER	<i>Perimeter of the mass</i>
AREA	<i>Area of the mass</i>
SMOOTHNESS	<i>Local variation in radius lengths</i>
COMPACTNESS	<i>Computed as: <math>\text{perimeter}^2 / \text{area} - 1.0</math></i>
CONCAVITY	<i>Severity of concave portions of the contour</i>
CONCAVE POINTS	<i>Number of concave portions of the contour</i>
SYMMETRY	<i>A measure of the symmetry of the nuclei</i>
FRACTAL DIMENSION	<i>'Coastline approximation' - 1.0</i>
DIAGNOSIS (Target)	<i>Diagnosis of cell sample: malignant or benign</i>

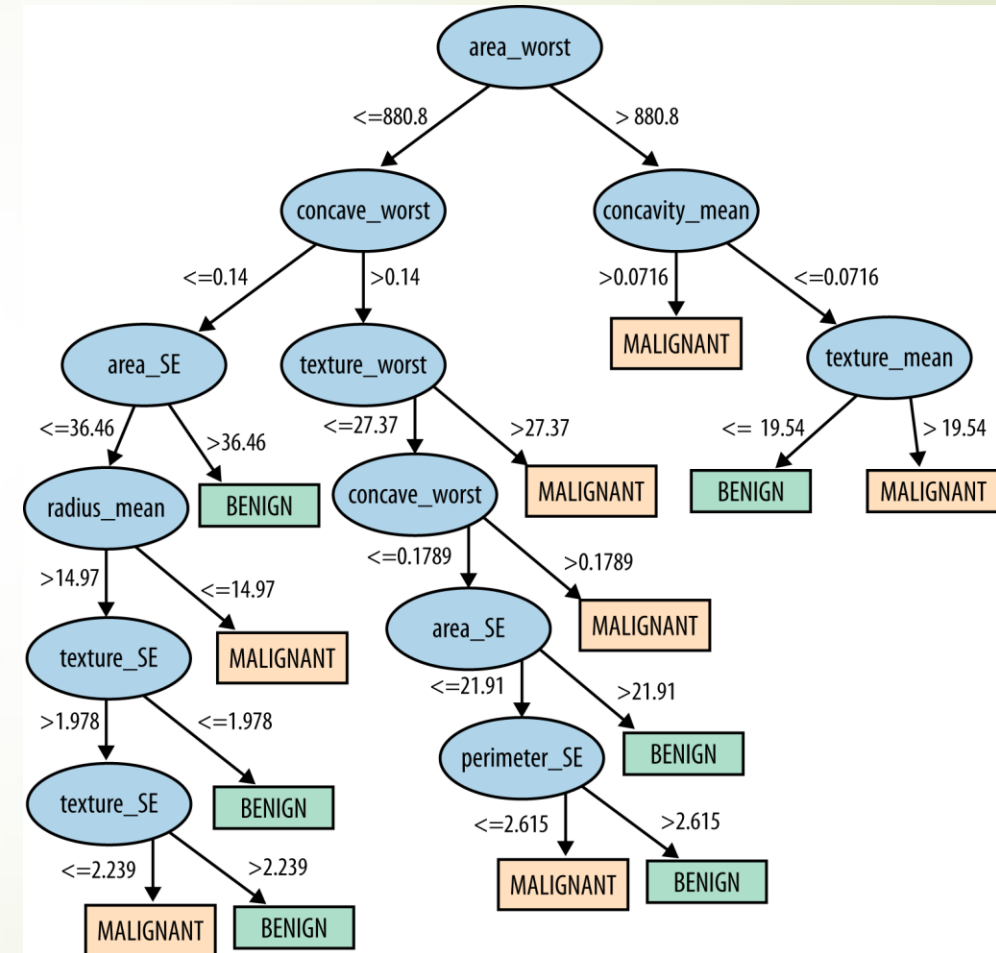
From each of these basic characteristics, three values were computed:

- the mean (`_mean`),
- standard error (`_SE`),
- “worst” or largest

# Logistic Regression vs. Decision Tree

Attribute	Weight (learned parameter)
SMOOTHNESS_worst	22.3
CONCAVE_mean	19.47
CONCAVE_worst	11.68
SYMMETRY_worst	4.99
CONCAVITY_worst	2.86
CONCAVITY_mean	2.34
RADIUS_worst	0.25
TEXTURE_worst	0.13
AREA_SE	0.06
TEXTURE_mean	0.03
TEXTURE_SE	-0.29
COMPACTNESS_mean	-7.1
COMPACTNESS_SE	-27.87
$w_0$ (intercept)	-17.7

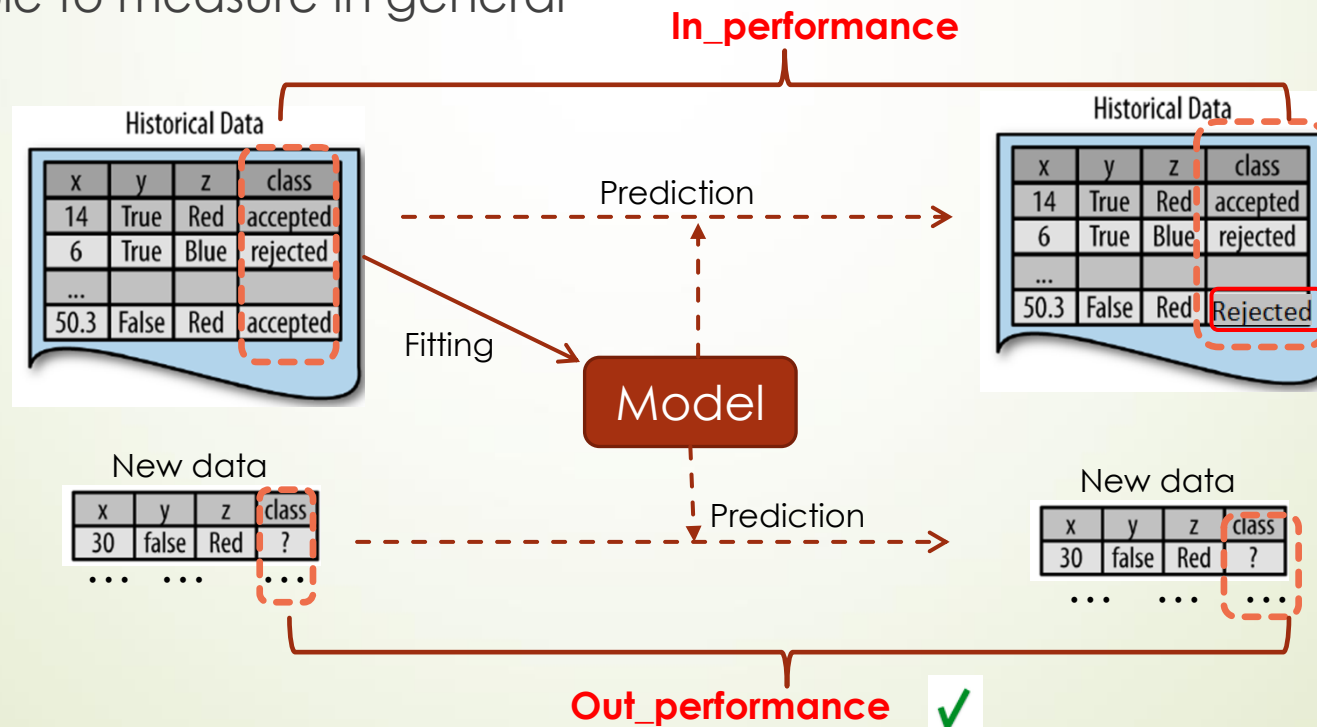
Linear equation learned by logistic regression  
acc. = 98.9%



Deducted decision tree(J48)  
acc. = 99.1%

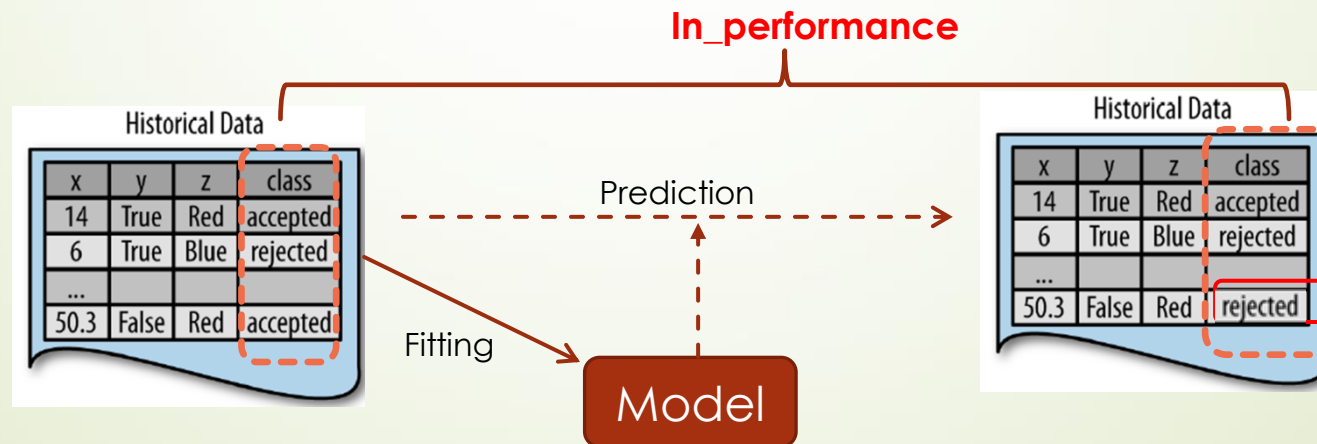
# Generalization

- We build predictive models to identify **general patterns** that predict the target attribute well for instances that we have not observed yet (Unknown new data)
- ✓ Good model: **generalize** well on the unknown data
- ✓ Model performance on the unknown data (**out performance**) is more important but impossible to measure in general



# Overfitting

- We try to optimize **in\_performance** when fitting predictive models to training data
  - ✓ We measure the in\_performance by error(/loss) or accuracy
  - ✓ But good in\_performance **DO NOT** guarantee good out\_performance
- Finding chance occurrences in data that look like interesting patterns, but which do not generalize, is called **overfitting the (training) data**
  - ✓ Overfitted model has **good in\_performance**, but **poor out\_performance**
  - ✓ We are easily fooled by overfitted model





# “Table” Model

- “Table” Model has perfect in\_performance (100% accuracy)
  - It does not make a single mistake, identifying correctly all the churners as well as the nonchurners.
  - Fitting: store the feature vector for each customer who has churned in a database table  $T_c$
  - Prediction: it takes the customer's feature vector, looks his/her up in  $T_c$ , and reports “100% likelihood of churning” if she is in  $T_c$  and “0% likelihood of churning” if he/she is not in
- “Table” Model but terrible out\_performance
  - Will predict “0% likelihood of churning” for every customer (not in the training data)
- A model that **looked perfect** would be completely **useless** in practice!



**Overfitting !**

# Overfitting and Model Complexity

- “If you torture the data long enough, it will confess” -- Ronald Coase
  - ✓ If we are allowed to use more complex models, we will have higher flexibility to pick up complex but effective patterns(good in\_performance)
  - ✓ More complex patterns are more likely to be just chance occurrences in the training data and then cannot generalize
  - ✓ Dubious patterns: second character of living city name for credit scoring ...
- Increase model complexity can improve in\_performance, but may not improve out\_performance (that's why overfitting)
  - ✓ Complexity of decision tree: nodes in the tree
  - ✓ Complexity of discriminant functions: dimension of feature vector, nonlinear or linear




# Trade-off of Model Complexity

- Trade-off between model complexity and the possibility of overfitting is fundamental in data mining
  - There is no single choice or silver bullet that will eliminate over-fitting, we need to recognize over-fitting and manage complexity in a principled way
  - All data mining procedures have **the tendency to over-fit** to some extent since they usually tailor models to the training data (some ones are better)




# Overfitting Mathematical Functions

- Increase the complexity of mathematical functions
  - ✓ Add more variables (more attributes):  $f(x) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5$
  - ✓ Extend existing variables:  $f(x) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_1^2 + w_7 * x_2/x_3$
- As you increase the dimensionality, you can perfectly fit larger and larger sets of arbitrary points
  - ✓ Carefully prune the attributes in order to avoid overfitting -> manual/auto selection





Classification



Regression


# Overfitting Decision Tree

- **Tree induction:** find important, predictive individual attributes recursively to smaller and smaller data subsets
  - ✓ We can build a perfect model if we are allowed to split the data/data subsets as many times as we can such that each leaf node is pure
  - ✓ The extreme case would be that each leaf node just contains one training example
  - ✓ Similar to look-up table, have perfect in\_performance but poor out\_performance (slightly better than look-up table)




# Examine Overfitting

- Overfitted model has good **in\_performance**, but poor **out\_performance**
- *Generalization performance* is measured by **out\_performance**
- It is hard to measure **out\_performance** on unknown data in general way
  - ✓ Approximating the out\_performance by performance evaluation conducted on data that have not been used for training (holdout data)
  - ✓ Split the whole training data into two parts: subset of training data and holdout data
  - ✓ Train model on training data, and then predict & evaluate performance on training data(training error) and holdout data (holdout error) respectively.



# Fitting Graph


- ▶ A fitting graph shows the performance(error/accuracy) of a model on training data and holdout data as a function of complexity
  - ✓ Focusing on the performance on the holdout data





# Fitting Graph of Table Model


- When new data comes to a table, make the prediction to be not churned
  - $N$  is the number of churned examples in training set
  - $M$  is total number of examples in training set,  $b=N/M$







# Fitting Graph of Decision Tree

- Measure model complexity by number of nodes in decision tree
  - ✓ Sweet spot represents the best trade-off between the extremes of (i) **not splitting** the data at all and simply using the average target value in the entire dataset, and (ii) building a complete tree out until **all** the leaves are pure.




# Drawbacks of Holdout Evaluation

- ▶ While a holdout set will indeed give us an estimate of generalization performance, it is just a single estimate.
- ✓ A single estimate of model accuracy might have just been a single particularly lucky (or unlucky) choice of training and test data
- ✓ Should we have any confidence in our estimation?




# How to Avoid Overfitting?


- Logistic regression are more likely to be overfitted than SVM




Original data



Adding one outlier



Adding one outlier



Adding one outlier with nonlinear feature

# Classification Models

## ► Linear classifier

- ✓ Logistic Regression: [https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)
- ✓ LinearSVC <https://scikit-learn.org/stable/modules/generated/sklearn.svm.LinearSVC.html>
- ✓ Perceptron [https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.Perceptron.html#sklearn.linear\\_model.Perceptron](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Perceptron.html#sklearn.linear_model.Perceptron)

## ► Non-linear classifier

- ✓ SVC: SVM for classification : <https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC>
- ✓ Decision Tree: <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html#sklearn.tree.DecisionTreeClassifier>

## ► Comparison

- ✓ [https://scikit-learn.org/stable/auto\\_examples/classification/plot\\_classifier\\_comparison.html#sphx-glr-auto-examples-classification-plot-classifier-comparison-py](https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html#sphx-glr-auto-examples-classification-plot-classifier-comparison-py)

# Outline

- Hinge Loss & SVM
- Nonlinear Models & Other Extensions
- Overfitting & Holdout Evaluation
- Quiz

# Lab Quiz

- **Deadline:** 17:59 p.m., Mar. 20, 2020
- Two questions accounting for **5%** of overall score
- **Upload** the **answer worksheet** and the accomplished **Python files** to the **Blackboard**
- You may submit **unlimited times** but only the **LAST** submission will be considered
- **Only the answers in answer sheet** will be referred for grading
- Note: **MUST attach ALL** the required files in every submission/resubmission, otherwise other files will be missing.